

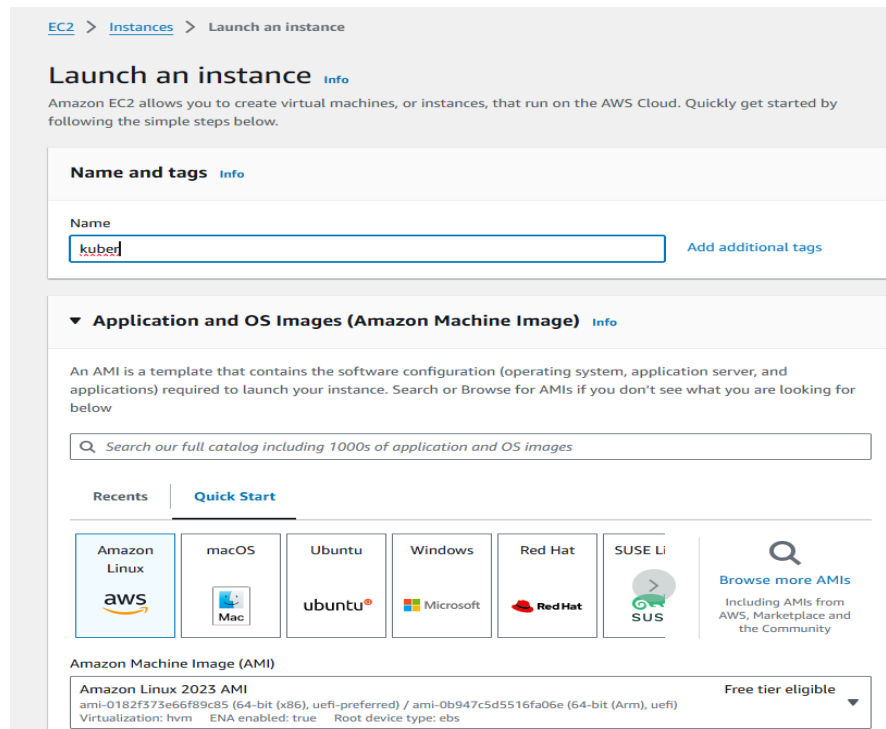
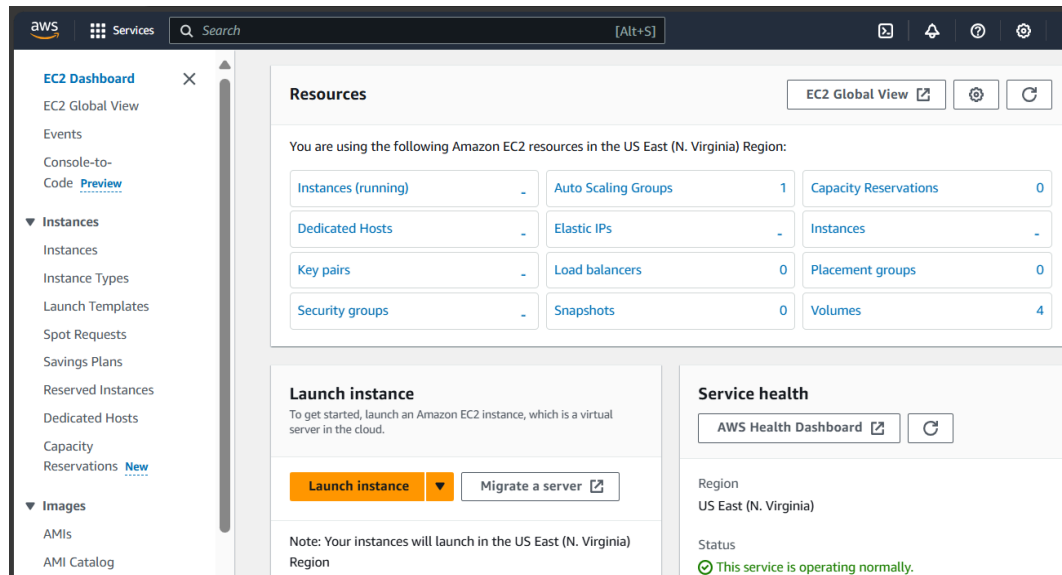
Experiment 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Procedure:

1. Creation Of EC-2 instance

- Create an EC2 AWS Linux instance on AWS .also edit the Security Group Inbound Rules to allow SSH. then select the t2.micro instance type



Description

Amazon Linux 2023 is a modern, general purpose Linux-based OS that comes with 5 years of long term support. It is optimized for AWS and designed to provide a secure, stable and high-performance execution environment to develop and run your cloud applications.

Architecture
64-bit (x86)

Boot mode
uefi-preferred

AMI ID
ami-0182f373e66f89c85

Verified provider

▼ Instance type
Info | Get advice

Instance type
t2.micro
Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand RHEL base pricing: 0.026 USD per Hour
On-Demand Linux base pricing: 0.0116 USD per Hour
Free tier eligible

All generations
Compare instance types

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login)
Info

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

server

Create new key pair

▼ Network settings
Info

Network
Info

vpc-051bba342b3626898

Subnet
Info

No preference (Default subnet in any availability zone)

Auto-assign public IP
Info

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)
Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group
Select existing security group

Common security groups
Info

Select security groups

test sg-0927cb1593f6ad5b2
VPC: vpc-051bba342b3626898

Compare security group rules

Security groups that you add or remove here will be added to or removed from all your network interfaces.

▼ Configure storage
Info

1x
8
GIB
gp3
Root volume (Not encrypted)

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage

Add new volume

Click refresh to view backup information

The tags that you assign determine whether the instance will be backed up by any

▼ Summary

Number of instances
Info

1

Software Image (AMI)

Amazon Linux 2023 AMI 2023.5.2...read more
ami-0182f373e66f89c85

Virtual server type (instance type)

t2.micro

Firewall (security group)

test

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million IOs, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel
Launch instance
Review commands

Instances (1)
Info

Find Instance by attribute or tag (case-sensitive)
All states

Instance ID = i-09dbca91fa3edcaea
Clear filters

Last updated less than a minute ago
Connect
Instance state
Actions
Launch instances

Name
Instance ID
Instance state
Instance type
Status check
Alarm status
Availability Zone
Public IPv4 DNS
Public IPv4 ...
Elastic IP

kuber
i-09dbca91fa3edcaea
Running
t2.micro
2/2 checks passed
us-east-1a
ec2-54-211-131-109.co...
54.211.131.109

- Thus Kuber named -instance gets created. Then click on Id of that instance then click on connect button you will see this:

2. Installation of Docker

1. . For installation of Docker into the machines run the following command:
`sudo yum install docker -y`

```
[ec2-user@ip-172-31-26-174 ~]$ sudo yum install docker -y
Last metadata expiration check: 0:05:13 ago on Fri Sep 13 13:17:25 2024.
Dependencies resolved.
=====
Package                                Architecture      Version            Repository
=====
Installing:
docker                                x86_64            25.0.6-1.amzn2023.0.2  amazonlinux
Installing dependencies:
containerd                            x86_64            1.7.20-1.amzn2023.0.1  amazonlinux
iptables-libs                         x86_64            1.8.8-3.amzn2023.0.2  amazonlinux
iptables-nft                         x86_64            1.8.8-3.amzn2023.0.2  amazonlinux
libcgroup                             x86_64            3.0-1.amzn2023.0.1    amazonlinux
libnetfilter_conntrack               x86_64            1.0.8-2.amzn2023.0.2  amazonlinux
libnftnl                             x86_64            1.0.1-19.amzn2023.0.2  amazonlinux
libnftnl                             x86_64            1.2.2-2.amzn2023.0.2  amazonlinux
pigz                                  x86_64            2.5-1.amzn2023.0.3    amazonlinux
runc                                  x86_64            1.1.13-1.amzn2023.0.1  amazonlinux
=====
Transaction Summary
=====
Total                                                                    78 MB/s | 84 MB    00:01
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing :
  Installing : runc-1.1.13-1.amzn2023.0.1.x86_64
  Installing : containerd-1.7.20-1.amzn2023.0.1.x86_64
  Running scriptlet: containerd-1.7.20-1.amzn2023.0.1.x86_64
  Installing : pigz-2.5-1.amzn2023.0.3.x86_64
  Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
  Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
  Installing : libnftnl-1.0.1-19.amzn2023.0.2.x86_64
  Installing : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
  Installing : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
  Installing : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
  Running scriptlet: iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
  Installing : libcgroup-3.0-1.amzn2023.0.1.x86_64
  Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64
  Installing : docker-25.0.6-1.amzn2023.0.2.x86_64
  Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.

  Verifying : containerd-1.7.20-1.amzn2023.0.1.x86_64
  Verifying : docker-25.0.6-1.amzn2023.0.2.x86_64
  Verifying : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
  Verifying : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
  Verifying : libcgroup-3.0-1.amzn2023.0.1.x86_64
  Verifying : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
  Verifying : libnftnl-1.0.1-19.amzn2023.0.2.x86_64
  Verifying : libnftnl-1.2.2-2.amzn2023.0.2.x86_64
  Verifying : pigz-2.5-1.amzn2023.0.3.x86_64
  Verifying : runc-1.1.13-1.amzn2023.0.1.x86_64

Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64    docker-25.0.6-1.amzn2023.0.2.x86_64    iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64    libcgroup-3.0-1.amzn2023.0.1.x86_64    libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnftnl-1.0.1-19.amzn2023.0.2.x86_64    libnftnl-1.2.2-2.amzn2023.0.2.x86_64    pigz-2.5-1.amzn2023.0.3.x86_64
runc-1.1.13-1.amzn2023.0.1.x86_64
Complete!
```

- Then, configure cgroup in a daemon.json file by using following commands

```
cd /etc/docker
```

```
cat <<EOF | sudo tee /etc/docker/daemon.json
```

```
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-26-174 ~]$ cd /etc/docker
[ec2-user@ip-172-31-26-174 docker]$ cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
```

- Then after this run the following command to enable and start docker and also to load the daemon.json file.

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-26-174 docker]$ sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-26-174 docker]$ sudo systemctl daemon-reload
[ec2-user@ip-172-31-26-174 docker]$ sudo systemctl restart docker
[ec2-user@ip-172-31-26-174 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

- docker -v

```
[ec2-user@ip-172-31-80-126 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

3. Then Install Kubernetes with the following command.

- SELinux needs to be disabled before configuring kubelet thus run the following command
sudo setenforce 0

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
[ec2-user@ip-172-31-26-174 docker]$ sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-26-174 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
```

- Here We are adding kubernetes using the repository whose command is given below.

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

```
[ec2-user@ip-172-31-26-174 docker]$ sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-26-174 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

- After that Run following command to make the updation and also to install kubelet ,kubeadm, kubectl: `sudo yum update`

```
[ec2-user@ip-172-31-80-126 docker]$ sudo yum update
Kubernetes
Dependencies resolved.
Nothing to do.
Complete!
100 kB/s | 17 kB    00:00
```

`sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes`

```
[ec2-user@ip-172-31-80-126 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Last metadata expiration check: 0:00:10 ago on Fri Sep 13 10:31:17 2024.
Dependencies resolved.
=====
Package                                Architecture      Version           Repository        Size
=====
Installing:
kubelet                                x86_64            1.30.5-150500.1.1  kubernetes        10 M
kubeadm                               x86_64            1.30.5-150500.1.1  kubernetes        10 M
kubectl                               x86_64            1.30.5-150500.1.1  kubernetes        17 M
Installing dependencies:
conntrack-tools                       x86_64            1.4.6-2.amzn2023.0.2  amazonlinux      208 k
cri-tools                             x86_64            1.30.1-150500.1.1    kubernetes        8.6 M
kubernetes-cni                        x86_64            1.4.0-150500.1.1     kubernetes        6.7 M
libnetfilter_cthelper                 x86_64            1.0.0-21.amzn2023.0.2  amazonlinux       24 k
libnetfilter_cttimeout                x86_64            1.0.0-19.amzn2023.0.2  amazonlinux       24 k
libnetfilter_queue                    x86_64            1.0.5-2.amzn2023.0.2  amazonlinux       30 k
=====
Transaction Summary
=====
Install  9 Packages
```

```
Total
Kubernetes
Importing GPG key 0x9A296436:
  Userid : "isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
  Fingerprint: E15 B144 86CD 377B 9E87 6E1A 2346 54DA 9A29 6436
  From : https://pkgs.k8s.io/core:/stable:/v1.30/rpm/repodata/repomd.xml.key
Key imported successfully
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing      : 1/1
  Installing     : kubernetes-cni-1.4.0-150500.1.1.x86_64 1/9
  Installing     : cri-tools-1.30.1-150500.1.1.x86_64 1/9
  Installing     : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 2/9
  Installing     : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 3/9
  Installing     : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 4/9
  Installing     : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 5/9
  Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 6/9
  Installing     : kubelet-1.30.5-150500.1.1.x86_64 7/9
  Running scriptlet: kubelet-1.30.5-150500.1.1.x86_64 7/9
  Installing     : kubeadm-1.30.5-150500.1.1.x86_64 8/9
  Installing     : kubectl-1.30.5-150500.1.1.x86_64 9/9
  Running scriptlet: kubectl-1.30.5-150500.1.1.x86_64 9/9
  Verifying      : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 1/9
  Verifying      : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 2/9
  Verifying      : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 3/9
  Verifying      : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 4/9
  Verifying      : cri-tools-1.30.1-150500.1.1.x86_64 5/9
  Verifying      : kubeadm-1.30.5-150500.1.1.x86_64 6/9
  Verifying      : kubectl-1.30.5-150500.1.1.x86_64 7/9
  Verifying      : kubelet-1.30.5-150500.1.1.x86_64 8/9
  Verifying      : kubernetes-cni-1.4.0-150500.1.1.x86_64 9/9

Installed:
conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.30.1-150500.1.1.x86_64      kubeadm-1.30.5-150500.1.1.x86_64
kubectl-1.30.5-150500.1.1.x86_64                kubelet-1.30.5-150500.1.1.x86_64      kubernetes-cni-1.4.0-150500.1.1.x86_64
libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
```

- After installing Kubernetes, we need to configure internet options to allow bridging.
 1. `sudo swapoff -a`

2. `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`
 3. `sudo sysctl -p`
- After installing Kubernetes, we need to configure internet options to allow bridging.

```
sudo swapoff -a
```

```
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
```

```
sudo sysctl -p
```

```
[ec2-user@ip-172-31-26-174 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
```

4. Initialize the Kubecluster

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```
[ec2-user@ip-172-31-80-126 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
I0913 10:32:44.629146 26680 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.30
[init] Using Kubernetes version: v1.30.4
[preflight] Running pre-flight checks
```

```
Your Kubernetes control-plane has initialized successfully!
```

```
To start using your cluster, you need to run the following as a regular user:
```

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
Alternatively, if you are the root user, you can run:
```

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

```
You should now deploy a pod network to the cluster.
```

```
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/
```

```
Then you can join any number of worker nodes by running the following on each as root:
```

```
kubeadm join 172.31.26.174:6443 --token pv0yyi.xh1lqhclfjr50pt8 \
--discovery-token-ca-cert-hash sha256:8293b2f6d29de466bd859007f5adbcdb3a
e6b0c446ba09033d32a5846b3d434f
```

- copy the token and save for future use .
`kubeadm join 172.31.26.174:6443 --token pv0yyi.xh1lqhclfjr50pt8 \`
`--discovery-token-ca-cert-hash`
`sha256:8293b2f6d29de466bd859007f5adbcdb3a`
`e6b0c446ba09033d32a5846b3d434f`
- Copy the mkdir and chown commands from the top and execute them
`mkdir -p $HOME/.kube`
`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```
[ec2-user@ip-172-31-80-126 docker]$ %
[ec2-user@ip-172-31-80-126 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Then, add a common networking plugin called flannel as mentioned in the code.

```
kubectl apply -f
https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

5. Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply deployment using this following command:

```
kubectl apply -f
https://k8s.io/examples/pods/simple-pod.yaml
```

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
```

Then use `kubectl get nodes` to check whether the pod gets created or not.

```
[ec2-user@ip-172-31-26-174 docker]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     0/1     Pending   0           12s
```

To convert state from pending to running use following command:

`kubectl describe pod nginx` This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.


```
[ec2-user@ip-172-31-26-174 docker]$ kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
Service Account: default
Node:          <none>
Labels:        <none>
Annotations:   <none>
Status:        Pending
IP:            <none>
IPs:           <none>
Containers:
  nginx:
    Image:      nginx:1.14.2
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-k4lj6 (ro)
```

```
Conditions:
  Type          Status
  PodScheduled  False
Volumes:
  kube-api-access-k4lj6:
    Type:          Projected (a volume that contains injected data from m
multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 3
00s
                    node.kubernetes.io/unreachable:NoExecute op=Exists for
300s
Events:
  Type          Reason          Age   From          Message
  ----          -
Warning        FailedScheduling  7s    default-scheduler  0/1 nodes are available: 1 no
de(s) had untoleraed taint {node-role.kubernetes.io/control-plane: }. preemption:
0/1 nodes are available: 1 Preemption is not helpful for scheduling.
```

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl taint nodes --all node-role.kubernetes.io
/control-plane-
node/ip-172-31-26-174.ec2.internal untainted
```

6. Now check pod status is is running

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   1 (6s ago)  90s
```

7. Lastly, mention the port you want to host. Here i have used localhost 8081 then check it.

```
kubectl port-forward nginx 8081:80
```

```
[ec2-user@ip-172-31-26-174 ~]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

8. Verify your deployment

Open up a new terminal and ssh to your EC2 instance.

Then, use this curl command to check if the Nginx server is running.

```
curl --head http://127.0.0.1:8080
```

If the response is 200 OK and you can see the Nginx server name, your deployment was successful. We have successfully deployed our Nginx server on our EC2 instance.

Conclusion: Firstly I created an EC2 AWS Linux instance successfully, then installed docker and kubernetes successfully, then initialized kubernetes which given me token and chown and mkdir command. Then I execute mkdir and chown the command successfully. Then I installed a networking plugin called flannel successfully. Then I tried to deploy nginx which initially gave an error. Then I deployed (simple-pod.yml) nginx successfully and also checked by using the get pods command, then hosted it on localhost 8081 ie <http://localhost:8081> successfully.